

Applications of distribution modeling and MCMC methods to intention forecasting

Francisco Valente Castro

September 30, 2020

Dr. Jean-Bernard Hayet

1. What is intention forecasting?
2. Learning the goal distribution
3. Goal inference using MCMC and CVAE

What is intention forecasting?

What is intention forecasting on pedestrians?



Figure 1: UCY university data-set. Prediction of possible trajectories. White - Real future trajectory, Green - Observed past, Orange - Sample of possible goals, Red - Sample of possible trajectories.

What is intention forecasting on pedestrians?

- We want to predict the intention of an agent (e.g. a pedestrian, a car) using information of a video recording and misc. information of the scene (e.g. semantic information of the image).
- On pedestrians, we want predict a possible goal (place they want reach, the intention) and a possible trajectory they are going to take, knowing all the information up to the present.
- A vehicle navigation system could predict the future possible positions of a pedestrian and adjust its own trajectory to avoid collisions.

What is intention forecasting on pedestrians?

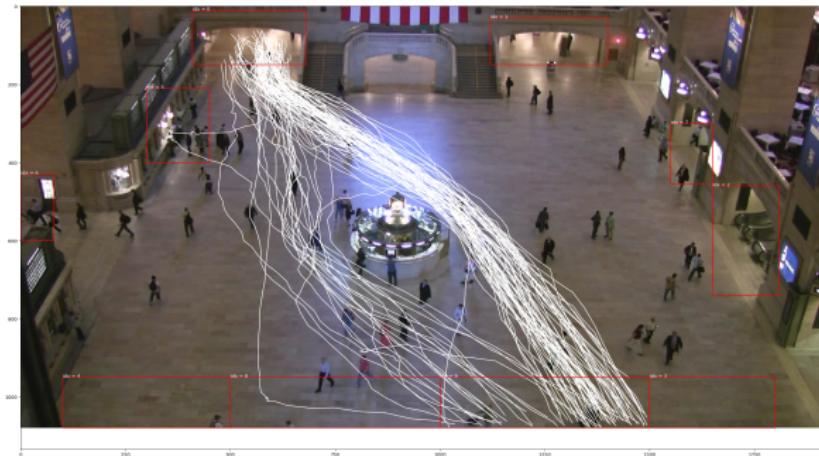


Figure 2: Grand Central Station pedestrian data-set exemplifying the data multi-modality.

Learning the goal distribution

Conditional Variational Auto-Encoder (CVAE)

The following are the random variables describing the principal components of our model

- $\mathbf{m} = (\mathbf{m}_1, \dots, \mathbf{m}_t)$: The intermediate points of the path that the pedestrian is taking.
- \mathbf{s} : The initial position of the trajectory.
- \mathbf{g} : The last position of the trajectory.
- \mathbf{z} : The variable in the latent space that describes the trajectory.

We want to train a generative model for the intermediate points \mathbf{m} in the trajectory, *conditioning* the response to the start (\mathbf{s}) and goal (\mathbf{g}).

What is intention forecasting on pedestrians?



Figure 3: Left: Samples from CVAE trained with a 'start' and 'goal' given as (x, y) coordinates. Right: Samples from CVAE trained with a 'start' and 'goal' given as polar angles to the edge of the scene.

Training the CVAE model

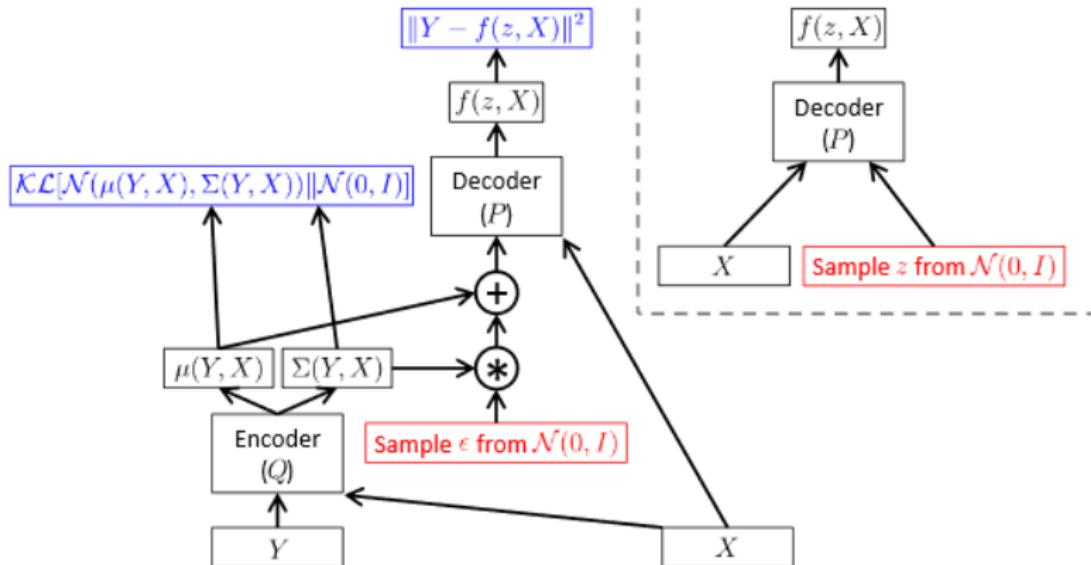


Figure 4: The general structure of a Conditional Variational Auto-encoder.

Training the CVAE model

Encoder

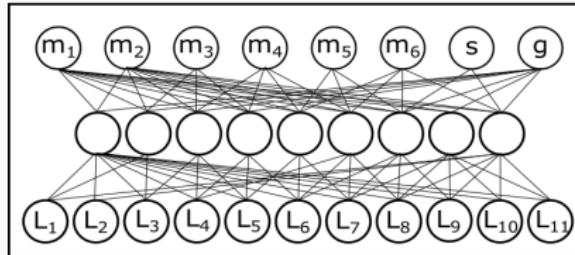


Figure 5: Linear Encoder $Q_\phi(\mathbf{m}, \mathbf{s}, \mathbf{g})$.

Decoder

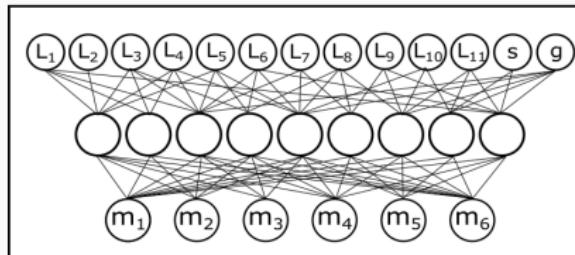


Figure 6: Linear Decoder $P_\theta(\mathbf{z}, \mathbf{s}, \mathbf{g})$.

Training the CVAE model

Minimize the ELBO cost function to train our Variational Auto-encoder.

$$\begin{aligned} \log P_\theta(\mathbf{m}|\mathbf{s}, \mathbf{g}) - D_{KL}[Q_\phi(\mathbf{z}|\mathbf{m}, \mathbf{s}, \mathbf{g})\|P(\mathbf{z}|\mathbf{m}, \mathbf{s}, \mathbf{g})] \\ = \mathbb{E}_{z \sim Q_\phi}[\log P_\theta(\mathbf{m}|\mathbf{z}, \mathbf{s}, \mathbf{g})] - D_{KL}[Q_\phi(\mathbf{z}|\mathbf{m}, \mathbf{s}, \mathbf{g})\|P(\mathbf{z}|\mathbf{s}, \mathbf{g})]. \end{aligned}$$

We are implicitly estimating the distributions $P_\theta(\mathbf{m}|\mathbf{z}, \mathbf{s}, \mathbf{g})$ (our *decoder*) y $Q_\phi(\mathbf{z}|\mathbf{m}, \mathbf{s}, \mathbf{g})$ (our *encoder*).

Goal inference using MCMC and CVAE

We find a way to estimate the posterior distribution of the **goal** given the **intermediate** points $\mathbf{m} = (\mathbf{m}_1, \dots, \mathbf{m}_t)$:

$$\begin{aligned}\mathbb{P}[\mathbf{g}|\mathbf{s}, \mathbf{m}] &= \frac{\mathbb{P}[\mathbf{g}, \mathbf{s}, \mathbf{m}]}{\mathbb{P}[\mathbf{s}, \mathbf{m}]}, \\ &= \frac{\mathbb{P}[\mathbf{m}|\mathbf{s}, \mathbf{g}]\mathbb{P}[\mathbf{s}, \mathbf{g}]}{\mathbb{P}[\mathbf{s}, \mathbf{m}]}, \\ &\propto \mathbb{P}[\mathbf{m}|\mathbf{s}, \mathbf{g}]\mathbb{P}[\mathbf{g}|\mathbf{s}]\end{aligned}$$

This **posterior** distribution $\mathbb{P}[\mathbf{g}|\mathbf{s}, \mathbf{m}] \propto \mathbb{P}[\mathbf{m}|\mathbf{s}, \mathbf{g}]\mathbb{P}[\mathbf{g}|\mathbf{s}]$ is the most important ingredient to apply MCMC methods to sample goals.

To actually calculate the probabilities in our posterior distribution $\mathbb{P}[\mathbf{g}|\mathbf{s}, \mathbf{m}] \propto \mathbb{P}[\mathbf{m}|\mathbf{s}, \mathbf{g}] \mathbb{P}[\mathbf{g}|\mathbf{s}]$ we have to sample μ means from our decoder.

$$\begin{aligned}\mathbb{P}[\mathbf{m}|\mathbf{s}, \mathbf{g}] &= \mathbb{E}_{\mathbf{z} \sim Q}[P(\mathbf{m}|\mathbf{z}, \mathbf{s}, \mathbf{g})], \\ &\propto \exp\left(-\frac{1}{2N} \sum_{i=1}^N (\mathbf{m} - \boldsymbol{\mu}_i)^T (\mathbf{m} - \boldsymbol{\mu}_i)\right).\end{aligned}$$

To estimate the required prior distribution $\mathbb{P}[\mathbf{g}|\mathbf{s}]$ we will use simpler methods like using **Gaussian Kernel Density Estimation**. We will show two different ways to do this estimation.

Prior 1D - Gaussian Kernel Density Estimation

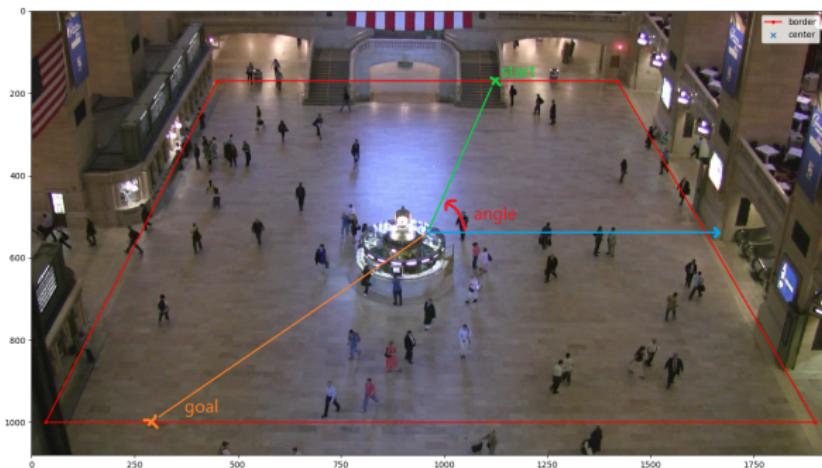


Figure 7: Start, Goal on border.

In this model we determine that **start** and **goal** are uni-dimensional random variables that take values on the range $(0, 2\pi)$ as angles.

Prior 2D - Gaussian Kernel Density Estimation

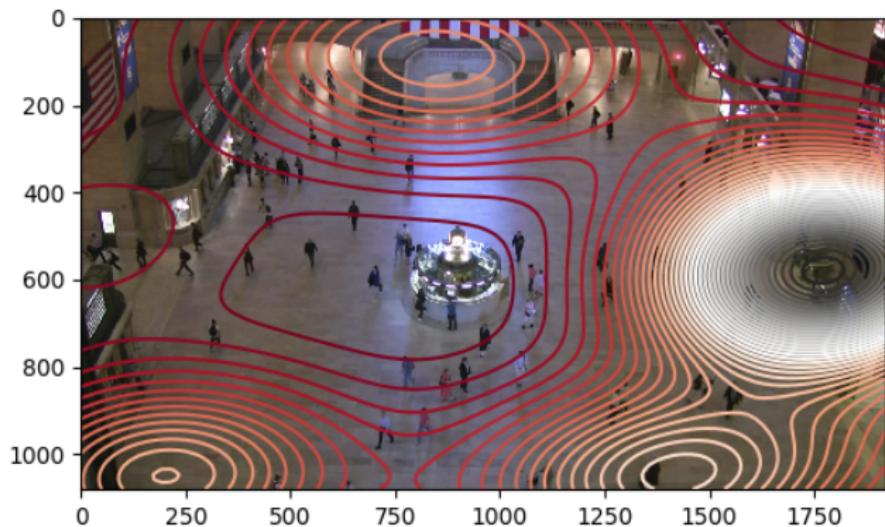


Figure 8: Gaussian Kernel Density Estimation of $\mathbb{P}[g]$.

2D Results on Grand Central dataset

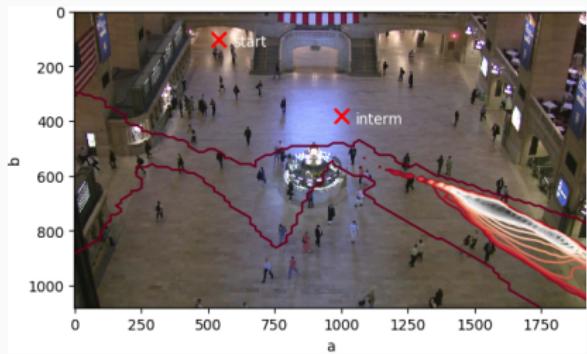
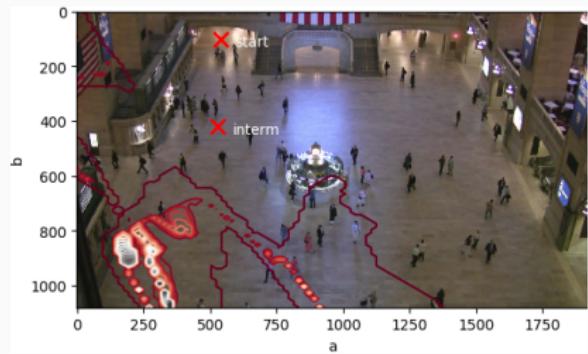


Figure 9: Posterior distribution estimation level sets.

2D Results on ETH dataset

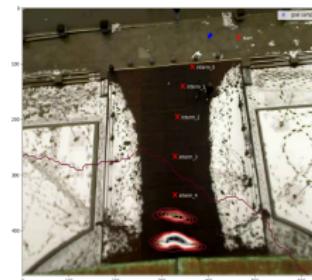
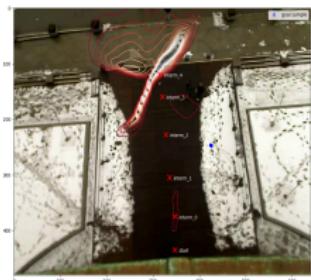
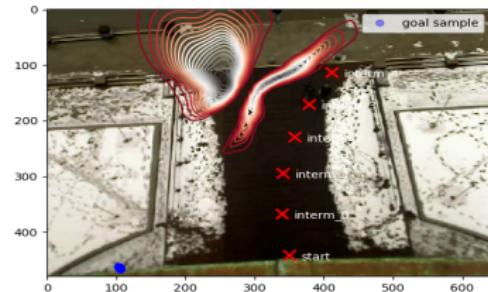
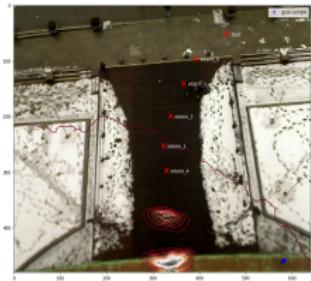


Figure 10: Posterior distribution estimation, ETH Dataset.

Choosing one or more proposals $q(\mathbf{g}^*|\mathbf{g}, \mathbf{s}, \mathbf{m}, \mathbf{M})$ (considering the algorithm using hybrid kernels) we can only consider the **acceptance probability** from the proposals in each iteration of the algorithm,

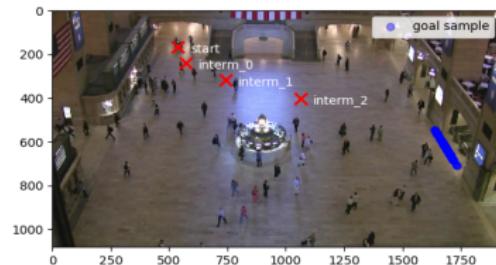
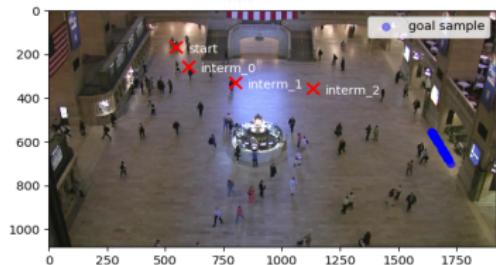
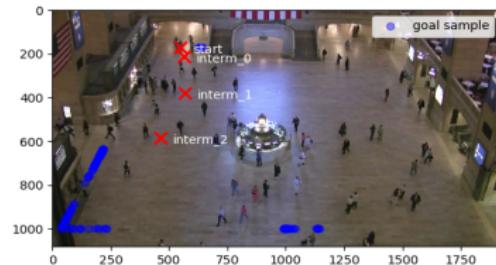
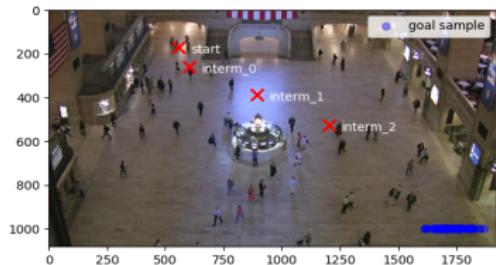
$$\alpha(\mathbf{g}^*|\mathbf{g}) = \min \left\{ 1, \frac{q(\mathbf{g}|\mathbf{g}^*)p(\mathbf{g}^*|\mathbf{s}, \mathbf{M})p(\mathbf{m}|\mathbf{g}^*, \mathbf{s}, \mathbf{M})}{q(\mathbf{g}^*|\mathbf{g})p(\mathbf{g}|\mathbf{s}, \mathbf{M})p(\mathbf{m}|\mathbf{g}, \mathbf{s}, \mathbf{M})} \right\},$$

In the following figure we can observe the **Metropolis-Hastings** algorithm in its simpler form. The variants and other MCMC methods follow a similar procedure.

Algorithm 1 Metropolis-Hastings algorithm

```
Initialize  $x^{(0)} \sim q(x)$ 
for iteration  $i = 1, 2, \dots$  do
    Propose:  $x^{cand} \sim q(x^{(i)} | x^{(i-1)})$ 
    Acceptance Probability:
         $\alpha(x^{cand} | x^{(i-1)}) = \min \{1, \frac{q(x^{(i-1)} | x^{cand})\pi(x^{cand})}{q(x^{cand} | x^{(i-1)})\pi(x^{(i-1)})}\}$ 
         $u \sim \text{Uniform}(u; 0, 1)$ 
        if  $u < \alpha$  then
            Accept the proposal:  $x^{(i)} \leftarrow x^{cand}$ 
        else
            Reject the proposal:  $x^{(i)} \leftarrow x^{(i-1)}$ 
        end if
    end for
```

Angle Results on GC using MCMC sampling



Thank you

References

- [1] Carl Doersch *Tutorial on Variational Autoencoders*. Carnegie Mellon / UC Berkeley, August 16, 2016
- [2] Ilker Yildirim *Bayesian Inference: Metropolis-Hastings Sampling*. Department of Brain and Cognitive SciencesUniversity of Rochester, 2012
- [3] Karttikeya Mangalam, Jitendra Malik *It Is Not the Journey but the Destination: Endpoint Conditioned Trajectory Prediction*. University of California, Berkeley